

Saket Deshmukh

GenAI / AI Engineer

Pune, Maharashtra | +91 7666114965 | saketdeshmukh10@gmail.com | linkedin.com/in/saket-deshmukh

Summary

GenAI Engineer with 2.5+ years of experience designing and shipping production-grade LLM applications, specializing in RAG pipelines, agentic workflows, and LLM evaluation. Skilled at building reliable AI systems with guardrails, grounding, and observability, while optimizing latency and inference costs. Experienced across the full stack: cloud infrastructure, containerization, and CI/CD for deploying AI applications in production.

Skills

GenAI: RAG, Agentic AI, LLM Applications, Prompt Engineering, Guardrails, Grounding

LLMs & LLMOps: LLM Evaluation, LLMOps, LangSmith, TruLens, OpenAI API

AI Frameworks: LangChain, LangGraph, FastAPI, REST APIs

Retrieval: FAISS, BM25, Hybrid Search, Embeddings, Reciprocal Rank Fusion (RRF)

Search: Semantic Search, Vector Search

Programming: Python, SQL, Bash

Cloud & Infrastructure: AWS, Azure, Docker, Kubernetes, Helm, Terraform, Ansible

CI/CD & DevOps: GitHub Actions, Jenkins, GitLab CI, Git, Linux

Observability: LangSmith, Prometheus, Grafana, AWS CloudWatch

Experience

KWAD Corporation Pvt. Ltd. – AI Engineer

AI Compliance & Financial Operations Copilot

Apr 2025 – Present

Pune, India

- Designed and built an agentic RAG system using LangChain and LangGraph for natural-language investigation of B2B transaction and compliance data, with multi-step workflows for conditional routing, retrieval, response generation, and human-in-the-loop escalation on complex queries.
- Implemented tenant-isolated retrieval, grounded responses with source citations, and input/output guardrails to ensure data separation, improve answer traceability, and prevent unsupported or unsafe requests from reaching downstream components.
- Deployed the application on Kubernetes using Docker and Helm with GitHub Actions CI/CD, and used LangSmith for tracing and RAG evaluation.

Briowin Private Limited – Software Engineer

CatalogSage – AI-Powered Product Catalog Assistant

Feb 2024 – Apr 2025

Mumbai, India

- Built an AI-powered RAG product assistant enabling natural-language search and recommendations across a 300+ SKU catalog using Python, FastAPI, LangGraph, and OpenAI.
- Designed a LangGraph workflow with conditional routing for input validation, hybrid retrieval, response generation, output validation, and graceful fallback.
- Engineered hybrid retrieval using OpenAI embeddings with FAISS semantic search and BM25 keyword retrieval, merging results using Reciprocal Rank Fusion (RRF).
- Built a RAG evaluation pipeline using TruLens to measure groundedness, context relevance, and answer relevance; containerized the service with Docker for deterministic CI testing.

Projects

Self-Hosted LLM Inference Platform

AWS EKS | Kubernetes | Terraform | Docker | vLLM

- Built and deployed a self-hosted LLM inference platform on AWS EKS, using Docker and vLLM for GPU-accelerated model serving, and Terraform to provision the underlying VPC, IAM, and Kubernetes infrastructure.

Education

B.E. in Computer Engineering

Savitribai Phule Pune University

2019 – 2023

Pune, India